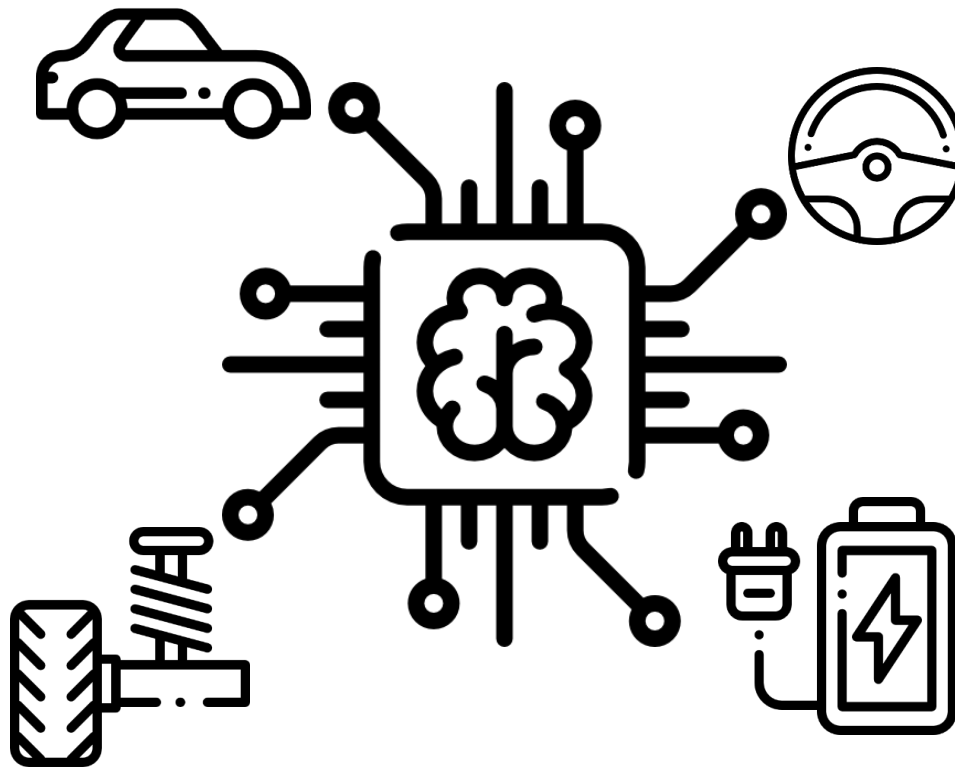


Artificial Intelligence in Automotive Technology

Maximilian Geißlinger / Fabian Netzler

Prof. Dr.-Ing. Markus Lienkamp





Lecture Overview

Lecture 16:15-17:45 Practice 17:45-18:30	
1 Introduction: Artificial Intelligence	20.10.2022 – Maximilian Geißlinger
2 Perception	27.10.2022 – Sebastian Huber
3 Supervised Learning: Regression	03.11.2022 – Fabian Netzler
4 Supervised Learning: Classification	10.11.2022 – Andreas Schimpe
5 Unsupervised Learning: Clustering	17.11.2022 – Andreas Schimpe
6 Introduction: Artificial Neural Networks	24.11.2022 – Lennart Adenaw
7 Deep Neural Networks	08.12.2022 – Domagoj Majstorovic
8 Convolutional Neural Networks	15.12.2022 – Domagoj Majstorovic
9 Knowledge Graphs	12.01.2023 – Fabian Netzler
10 Recurrent Neural Networks	19.01.2023 – Matthias Rowold
11 Reinforcement Learning	26.01.2023 – Levent Ögretmen
12 AI-Development	02.02.2023 – Maximilian Geißlinger
13 Guest Lecture	09.02.2023 – to be announced

Objectives for Lecture 4: Classification

After the lecture you are able to...

... understand the concept of classification, its association to pattern recognition and the urge for machine learning.

... acquire labeled training data and prepare it for the training and validation phase.

... plan the basic workflow for an arbitrary supervised learning problem.

... understand the concepts of different classification methods together with their pro and cons.

... implement, train and use a classification method with python libraries.

... understand how classification can be used in the perception for automated vehicles.

... analyze the quality of a given classifier regarding to different criteria.



Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

(Andreas Schimpe, M.Sc.)

Agenda

1. Chapter: Introduction

1.1 Overview

1.2 Training and Validation

2. Chapter: Methods



2.1 Logistic Regression

2.2 Nearest Neighbors

2.3 Support Vector Machine

3. Chapter: Application

4. Chapter: Summary



Classification

*“Systematic arrangement in groups or categories
according to established criteria” [13]*

[3]



Additional Slide

Classification, independent of machine learning and AI, is about assigning something perceived to a known class. The perceived can be an object, a process, a behavior or similar. The classes in our everyday life can be arbitrarily generic or specific, and can also overlap. If we blindly drink a beer we can classify it as "liquid", but we can also go much further and classify it as a "drink", a "beer" or "bitter". In dependence of our training (if we have often drunk a beer of the same brand), we can also classify the brand although it may not be written on the bottle.

A core element of classification is that we have to learn rules before we can classify. If someone, who has drunk wine all his life and has never heard or seen anything about beer, drinks the first beer, he will not be able to recognize it as beer. Only when he learns what characteristics a beer has (appearance, taste,...), he will be able to classify a drink as beer.

The next thing to emphasize is that classification is not about memorizing something, but about building an abstract model. When we drink the first beer, we remember properties (features) and build rules that have to be met in order to classify a drink as beer. This way we can later classify, what we perceived even though we had never perceived it before.

Classification

“Systematic arrangement in groups or categories according to established criteria” [13]

[14]




Method Overview

Pattern Recognition


Regression

- Predict **continuous** valued output
- Supervised




Classification

- Predict **discrete** valued output
- Supervised



Clustering

- Predict discrete valued output
- **Unsupervised**



Additional Slide

Patterns are distinctive structures in data. The data can have an arbitrary origin, so patterns can be searched in images, in behavior, in processes, in measurement signals etc., in principle everywhere, where data (or only metadata) is available.

The patterns can indicate a trend (regression).

The patterns can refer to a certain known type (classification).


The patterns can indicate similarities between different data points (clustering).

Pattern recognition in machine learning is based on the ability to generate abstract rules for patterns from a large amount of data. With the help of these rules, it is possible to gain information about the data structure (clustering), forecasts (regression) and to better categorize new data (classification).

It is important that many data have to be available in good quality for pattern recognition to work with machine learning


Method Overview – Examples

Regression




- House pricing
- Number of sales
- Persons weight

Classification



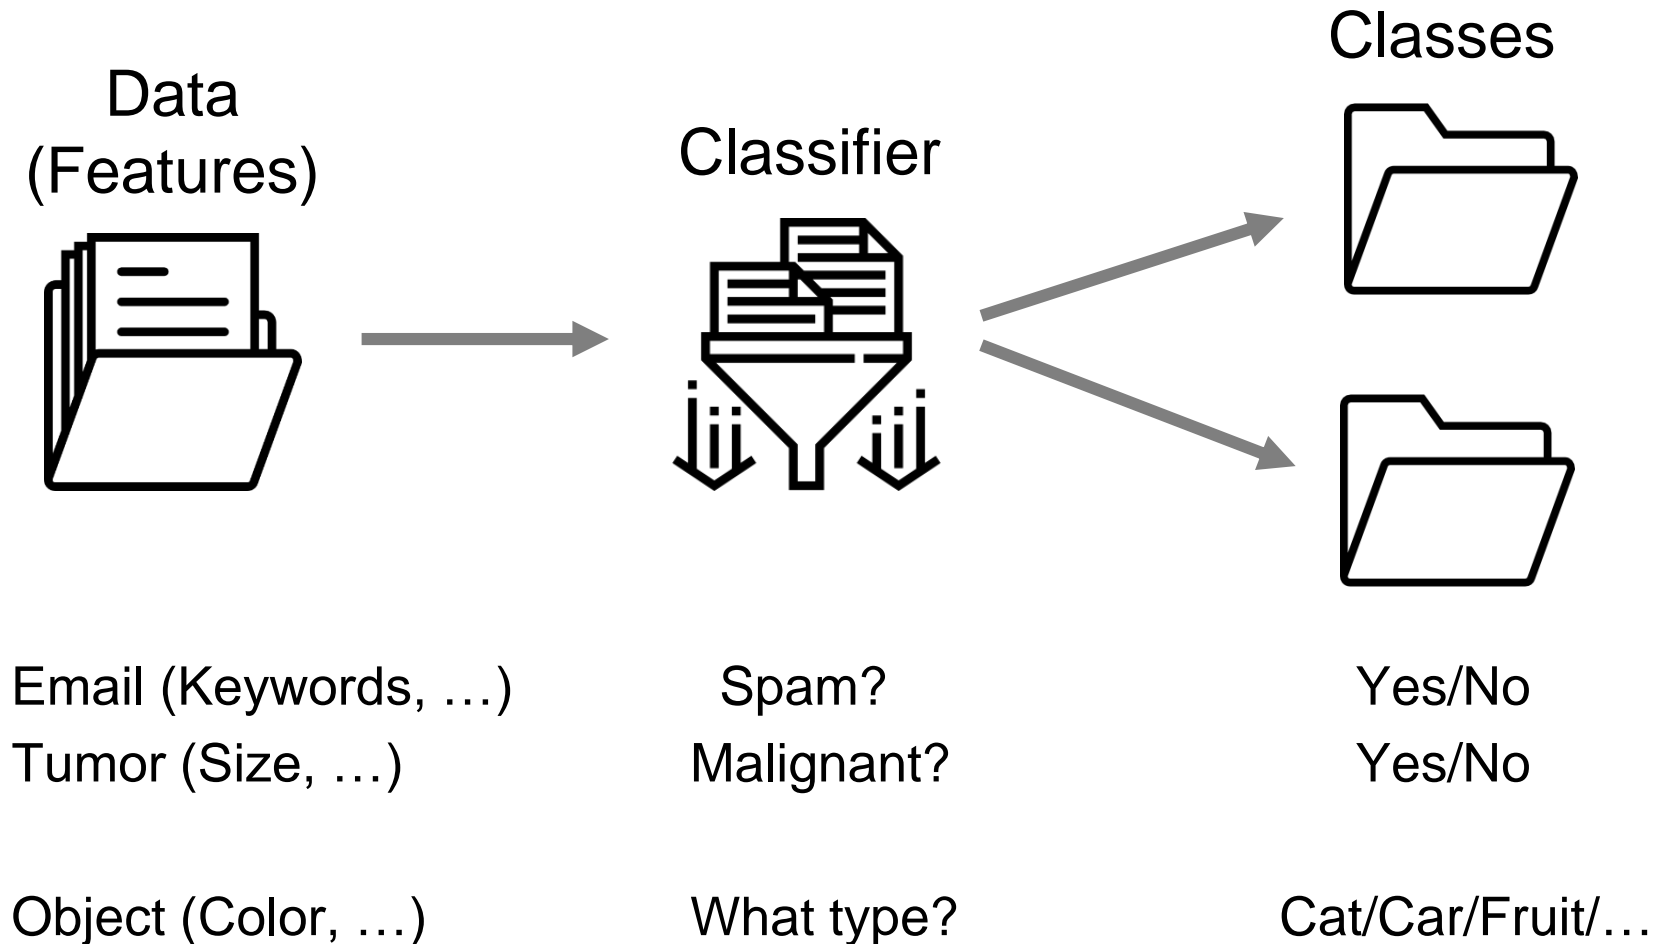
- Object detection
- Spam detection
- Cancer detection

Clustering

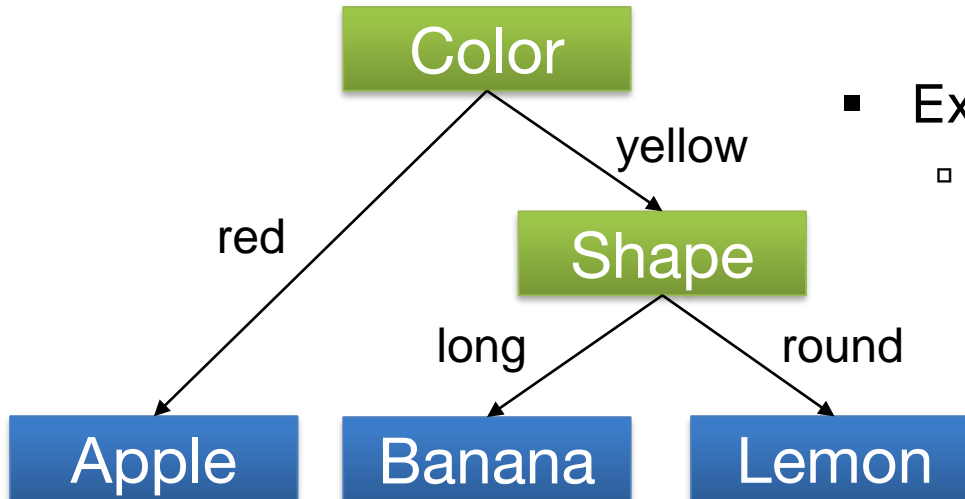


- Genome patterns
- Google news
- Pointcloud (Lidar) processing

General Approach



Classic Method vs. Machine Learning Method



- Example: Decision tree
 - Use a-priori knowledge to formulate classification rules


- Advantages of machine learning
 - Automatic generation of a-priori knowledge
 - Automatic generation of complex classification rules
- Suitable for extreme large datasets

Machine Learning in Reality



[16]

Classification – Example



Object
Classification



Object
Detection



Object
Tracking

[4]

Additional Slide

In this video, different vehicle types were detected. However, it is also noticeable that no vehicles are detected on the opposite lane. Two potential reasons for this could be that

- the image area that is classified is limited to the center area of the camera images
- or only vehicle fronts and no rear of vehicles can be detected as training data only included labels for vehicle fronts.


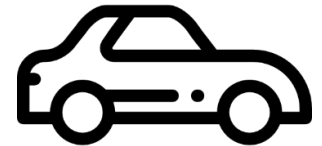
Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

(Andreas Schimpe, M.Sc.)

Agenda

1. Chapter: Introduction
 - 1.1 Overview
 - 1.2 Training and Validation**
2. Chapter: Methods
 - 2.1 Logistic Regression
 - 2.2 Nearest Neighbors
 - 2.3 Support Vector Machine
3. Chapter: Application
4. Chapter: Summary

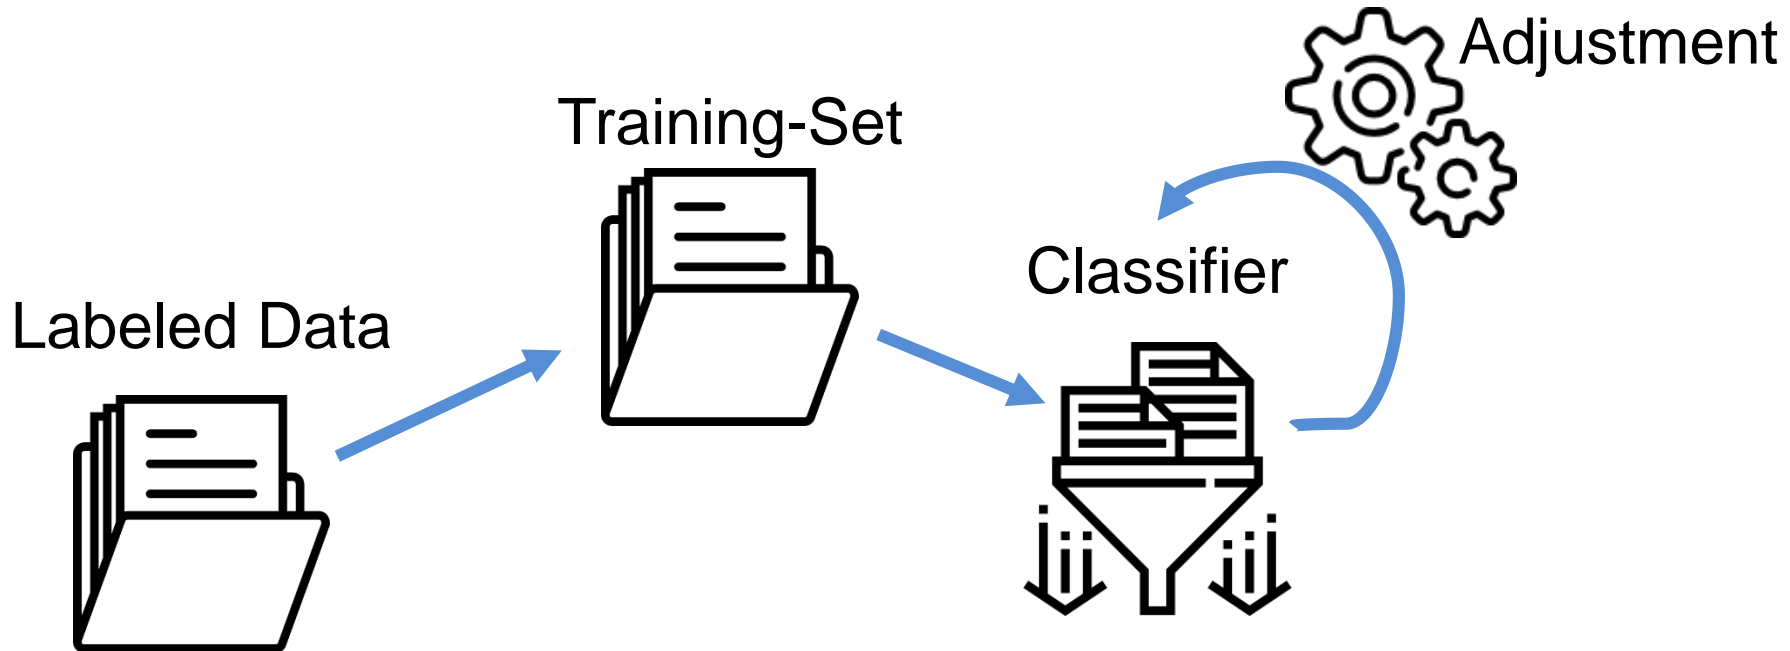


Formal Definition - Classification

$$C_{M(\theta)}: X \rightarrow Y$$


- **Classifier:** C
- **Model:** M with parameters θ
- **Input:** Feature Space $X = [x_1, x_2, \dots, x_N]^T$
- **Output:** (Predicted) Labels / Classes: Y
- **Dataset:** $D = \{(X_1, Y_1), (X_2, Y_2), \dots, (X_M, Y_M)\}$
- **Training:**
 - Given training data set $O \subseteq D$ with known labels
 - Optimize model parameters θ by minimizing prediction errors
- **Classification:**
 - Evaluate $C_{M(\theta)}$ at (unknown/new) element X
 - Receive class prediction Y

Supervised Learning - Classification



— Training

Classifier Training Result

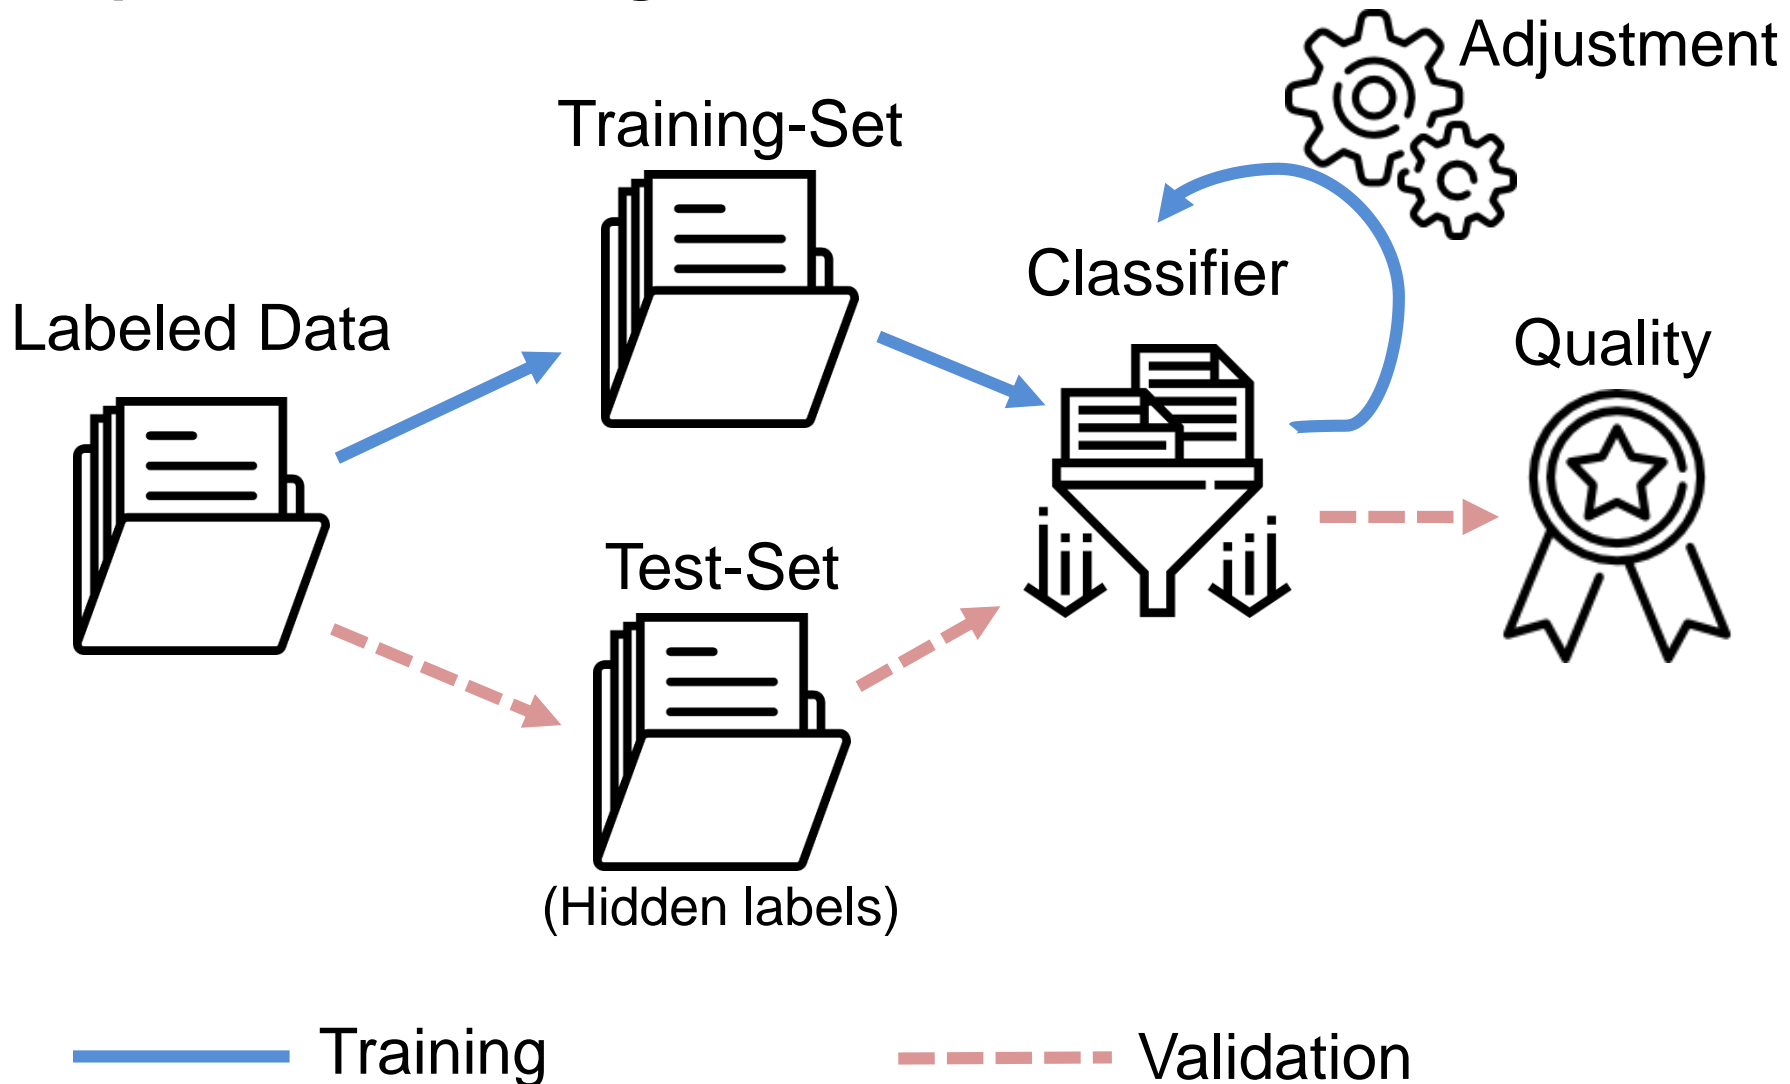


Features: Size, Weight

Data Points: ○

Decision Boundaries:

Supervised Learning - Classification



Additional Slide

For supervised learning, labelled data must be available. These are divided into a training set (approx. 80 %) and a test set (approx. 20 %). With the help of the training data, a machine learns, optimizing its model parameters. Afterwards, the (unseen) test data is used to check the quality of the model.

Quality Measures for Classifiers

Scalability
Compactness
Accuracy
Interpretability
Efficiency
Robustness

Quality Measures for Classifiers

- Classification accuracy or classification error (complementary)
 - Loss Functions
- Compactness of the model
 - Decision tree size; number of decision rules
- Interpretability of the model
 - Insights and understanding of the data provided by the model
- Efficiency
 - Time to generate the model (training time)
 - Time to apply the model (prediction time)
- Scalability for large databases
 - Efficiency in disk-resident databases
- Robustness
 - Robust against noise or missing values

Validation of Classifiers

- k-fold Cross Validation
 - Decompose data set evenly into k subsets of (nearly) equal size
 - Iteratively use k-1 partitions as training data and the remaining single partition as test data.



- Additional requirement: stratified folds
 - Class distributions in training and test set should represent the class distribution in whole data set (or at least in training set)
- Standard: 10-fold stratified cross validation

Confusion Matrix

		Classified as			
Correct Label		Class 1	Class 2	Class 3	Class 4
	Class 1	45 TP	0	2 FN	1
	Class 2	3	44	0	1
	Class 3	0 FP	0	67 TN	0
	Class 4	8	5	6	37

- Recall: $\frac{TP}{TP+FN}$

- Precision: $\frac{TP}{TP+FP}$

- Specificity: $\frac{TN}{TN+FP}$

True Positives: TP

False Positives: FP

True Negatives: TN

False Negatives: FN

Additional Slide

There is always a Trade Off between two values of FN TN TP TN

Recall - **Sensitivität**: Fraction of test objects of class X, which have been identified correctly

Precision - **Genauigkeit**: Fraction of test objects assigned to class X, which have been identified correctly

Specificity - **Spezifität**: Fraction of test objects assigned to other classes, which actually belong to other classes

Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

(Andreas Schimpe, M.Sc.)

Agenda

1. Chapter: Introduction

1.1 Overview

1.2 Training and Validation

2. Chapter: Methods



2.1 Logistic Regression

2.2 Nearest Neighbors

2.3 Support Vector Machine


3. Chapter: Application

4. Chapter: Summary



Methods

- Decision Trees
- **Logistic Regression**
- **Nearest Neighbors**
- **Support Vector Machine**
- Neural Networks
- ...



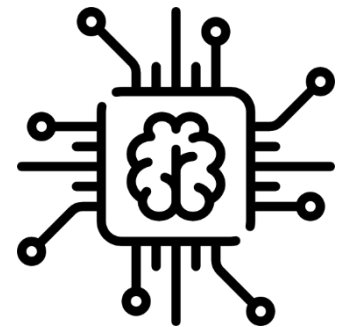
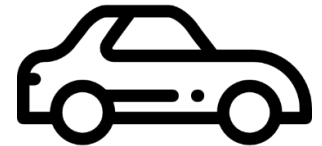
Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

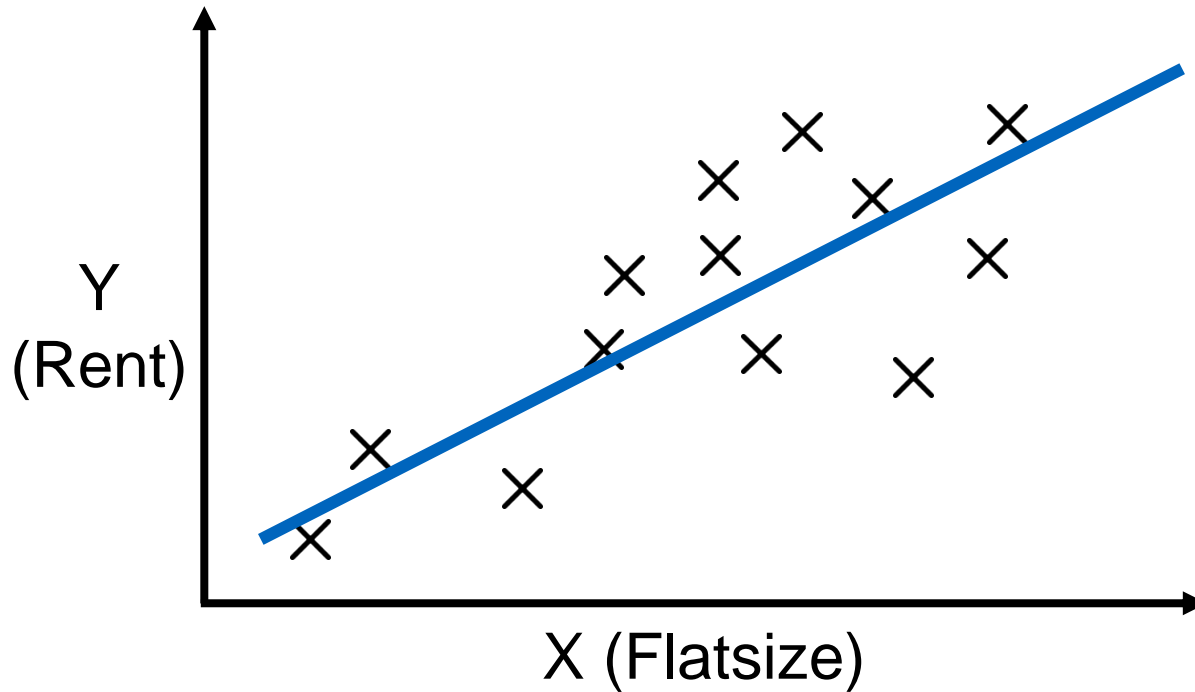
(Andreas Schimpe, M.Sc.)

Agenda

1. Chapter: Introduction
 - 1.1 Overview
 - 1.2 Training and Validation
2. Chapter: Methods
 - 2.1 Logistic Regression**
 - 2.2 Nearest Neighbors
 - 2.3 Support Vector Machine
3. Chapter: Application
4. Chapter: Summary




Recap Linear Regression



$$Y = h_{\theta}(X), \quad Y \in \mathbb{R}$$

Linear Regression for Classification




$$Y = h_{\theta}(X), \quad Y \in \mathbb{R}$$


Predict malignant if $Y \geq 0.5$

Sigmoid Function

$$g(z) = \frac{1}{1 + e^{-z}}$$



Logistic Regression



Probabilistic classification:

$$Y = g_{\theta}(h_{\theta}(X)) \quad Y \in]0,1[$$

Predict malignant if $Y \geq 0.5$

Discussion Logistic Regression

Pro

- **Implementation:** Easy to use
- **Probabilistic:** Probability of an object being in a certain class
- **Computation:** Quick training phase
- **Insights:** Produces understandable models

Contra

- **Linearity:** Hard to adopt to non-linear problems
- **Overfitting:** Training data has to be well chosen

Additional Slide

Linear regression alone is not suitable for classifying data because a continuous output must be mapped to discrete values. This blurs the class boundaries, the most relevant element for classification. By using the Sigmoid function, a S function limited between 0 and 1, the logistic regression is obtained. Despite the name, this is a classification method. The rather steep flank of the S function makes it possible to build up a clear class boundary.

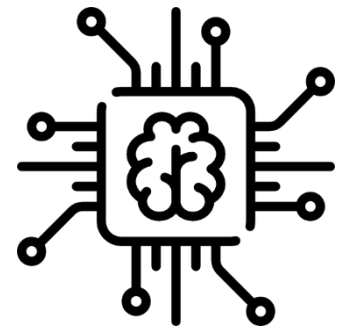
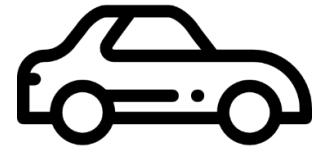
Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

(Andreas Schimpe, M.Sc.)


Agenda

1. Chapter: Introduction
 - 1.1 Overview
 - 1.2 Training and Validation
2. Chapter: Methods
 - 2.1 Logistic Regression
 - 2.2 Nearest Neighbors**
 - 2.3 Support Vector Machine
3. Chapter: Application
4. Chapter: Summary



Nearest Neighbors

Classify a new object based on its nearest neighbor(s)



Nearest Neighbors

An instance-based (memory-based) „learning“ method

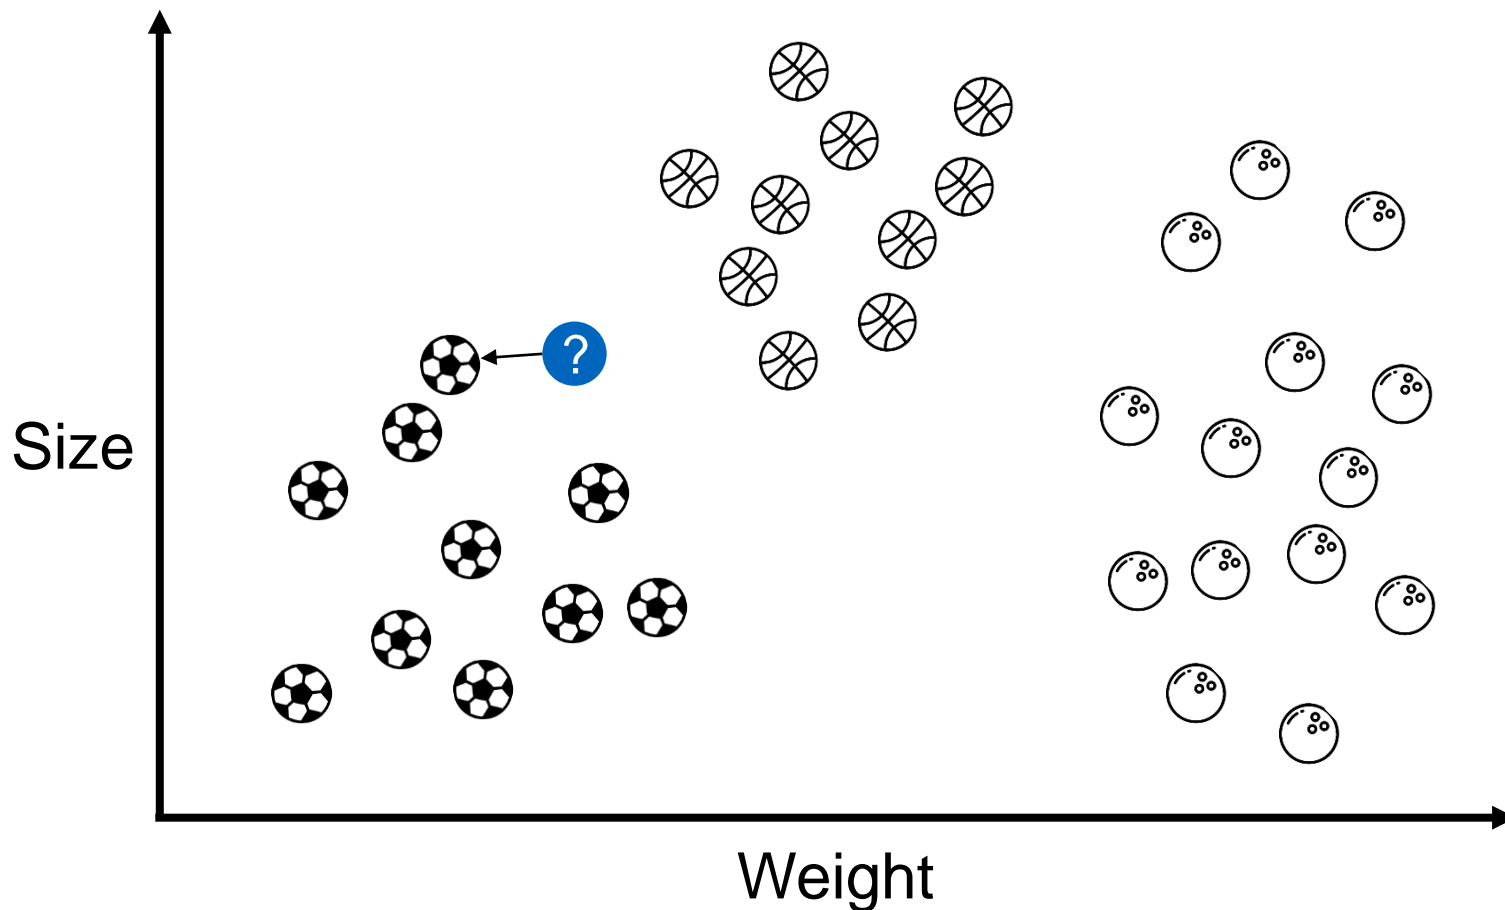
- No training and test phase → No model is generated
 - Instead: Store labeled training data as points in metric space
 - Process training data when new object should be classified
→ „lazy evaluation“
- Easy to use, but memory and time inefficient for large data sets

Nearest Neighbors Variants

- NN Classifier
 - Classify based on nearest neighbor only
- k-NN Classifier
 - Classify based on k nearest neighbors ($k > 1$)
- Weighted k-NN Classifier
 - Classify based on k nearest neighbors, weighted by their distances
- Mean-based NN Classifier
 - Classify based on distance to mean positions of classes


Nearest Neighbors Variants

- NN Classifier
 - Classify based on nearest neighbor only




Nearest Neighbors Variants

- k-NN Classifier
 - Classify based on k nearest neighbors ($k > 1$)




Nearest Neighbors Variants

- Weighted k-NN Classifier
 - Classify based on k nearest neighbors, weighted by their distances





Nearest Neighbors Variants

- Mean-based NN Classifier
 - Classify based on distance to mean positions of classes




k-NN Classifier

- How to choose k?
 - Generalization vs Overfitting
 - Large k: Many objects from different classes
 - Small k: Sensitivity against outliers
 - Practice: $1 \ll k < 10$



Weighted k-NN Classifier

- How to weight the neighbors?
 - Distance to the neighbor
 - $w_i = \frac{1}{distance_i^2}$
 - Frequency of the neighbors class
 - $w_i = \frac{1}{frequency_i}$




Classification Result

Normal ●

Weighted (Distance) ●

Weighted (Frequency) ●



Additional Slide

Nearest Neighbors classification does not build an explicit model compared to the other classification methods, but uses all existing data to check how the elements near the new data point look like for each new classification. Beside the direct proximity to the next element there are several other variants, especially k-Nearest Neighbors, where the k data points with the smallest distance to the new data point are considered. The choice of k is relevant, so a value must be found that balances between overfitting and generalization.

It is also possible to make a model from the Nearest Neighbors method. For $k = 1$, the results would be a Voronoi diagram. However in most cases, the effort to construct this is greater than the effort to compute the k nearest neighbors directly.

Discussion NN Classifier

Pro

- **Applicability:** Easy to calculate distances
- **Accuracy:** Great results for many applications
- **Incremental:** Easy adoption of new training data
- **Robust:** Scopes with noise by averaging (k-NN)

Contra

- **Efficiency:** Processing grows with training data $\mathcal{O}(n)$
 - Can be reduced to $\mathcal{O}(\log n)$ with an index structure (requires training phase)
- **Dimensionality:** Not every dimension is relevant
 - Weight dimensions (scale axes)

Neutral

- Does not produce explicit knowledge about classes



Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp


(Andreas Schimpe, M.Sc.)

Agenda

1. Chapter: Introduction
 - 1.1 Overview
 - 1.2 Training and Validation
2. Chapter: Methods
 - 2.1 Logistic Regression
 - 2.2 Nearest Neighbors
 - 2.3 Support Vector Machine**
3. Chapter: Application
4. Chapter: Summary




Support Vector Machine (SVM)




- Linear separation
 - Objects: $X \in \mathbb{R}^N$
 - Two classes: $Y \in \{-1, 1\}$
 - Hyperplane separates both classes
- Training
 - Compute hyperplane
- Classification
 - Compute sign / distance to hyperplane

SVM – Maximum Margin Hyperplane (MMH)



- **Objective**
Maximize distance to Hyperplane
 - Distance to hyperplane at least δ (Margin) for all data points
- High generalization
 - Maximal stable
- Small number of support vectors
 - Only depends on objects with distance δ

SVM – Formal Definition




- Trainingdata: $(X_1, Y_1) \dots (X_M, Y_M)$
with $X \in \mathbb{R}^N, Y \in \{-1, 1\}$
- Hyperplane: $w^T X + b = 0$
with $w \in \mathbb{R}^N$: normalvector,
 $\frac{b}{\|w\|}$: offset from origin
- Margin: $\delta = \frac{1}{\|w\|}$
- Training: Minimize $\|w\|$
subject to $Y_i(w^T X_i + b) \geq 1$, for $i = 1 \dots M$
- Classification of new sample X :
$$Y = \begin{cases} +1, & \text{if } w^T X + b \geq 0 \\ -1, & \text{else} \end{cases}$$


SVM – Soft Margin

Linear Separation

Not always optimal




Not always possible




- Tradeoff between error and margin
- Allow classification error to maximize margin (soft margin)

SVM – Space Transformation




- Non-linear data
→ Too many errors with soft margin

SVM – Space Transformation




- Non-linear data
→ Too many errors with soft margin
- Transform in higher dimensional space $\Phi(X)$
- Separation with linear hyperplane
 $w^T \Phi(X) + b = 0$
- Inverse transform yields non-linear hyperplane in original space
- Example: Quadratic transformation
 - Hyperplane becomes polynomial of degree 2

SVM – Space Transformation Visualisation



[5]

SVM – Space Transformations



1. Transform Data Points

- Lower to higher dimension
- Computational complex

2. SVM Training → Get Hyperplane

3. Transform Hyperplane

- Higher to lower dimension
- Computational complex

SVM – Kernel Machines

- SVM training and classification:
both functions of dot product of samples
- Kernel Trick: Replace dot product with non-linear kernel function
→ **Space transformation no longer required**
- Kernel functions:
 - Polynomial: $k(X_i, X_j) = (X_i^T X_j)^d$
 - Gaussian radial basis function (RBF):
 $k(X_i, X_j) = \exp(-\gamma \|X_i - X_j\|^2)$ for $\gamma > 0$
 - Linear, sigmoid, tanh, ...

SVM – Kernel Machines


Example: Polynomial space transformation (2nd degree)

- Original space (example: 2D): $X = [x_1, x_2] \in \mathbb{R}^2$
- Transformed space: $\Phi(X) = [x_1^2, x_1 x_2, x_2 x_1, x_2^2] \in \mathbb{R}^4$
- Scalar product of two samples in transformed space
$$\begin{aligned}\Phi(X_1)^T \Phi(X_2) &= x_{1,1}^2 x_{2,1}^2 + 2x_{1,1} x_{2,1} x_{1,2} x_{2,2} + x_{1,2}^2 x_{2,2}^2 \\ &= (x_{1,1} x_{2,1} + x_{1,2} x_{2,2})^2\end{aligned}$$
- Polynomial kernel function (example: 2nd degree)
$$k(X_1, X_2) = (X_1^T X_2)^2 = (x_{1,1} x_{2,1} + x_{1,2} x_{2,2})^2 = \Phi(X_1)^T \Phi(X_2)$$


→ No transform in higher dimensional space required

SVM – Kernel Machines


SVC with linear kernel




LinearSVC (linear kernel)



SVC with RBF kernel



SVC with polynomial (degree 3) kernel




[15]


Multi-Class SVM

Combination of SVMs

1 vs. Rest



1 vs. 1



Additional Slide

SVM are very powerful classifiers that are still widely used today. The basic idea is to separate the classes by a clear boundary. The border should have the maximum distance to both classes, i.e. run through the middle of the lane between the classes. The points of the different classes closest to the class boundaries are the support vectors and "span" the class boundary. The naive idea is that, for example, in 2D a straight line can be drawn between the classes, and in 3D a plane separates the classes. This is often not the case, so a curved line or plane would be necessary. The approach of SVM is to transform the data points into a higher dimensional space until a linear separation by a hyperplane is possible. Since a transformation into a higher space is very computationally expensive, the kernel trick is used, which allows to calculate relations between two data points in a highly transformed space efficiently. Therewith, the whole transformation into a higher dimension can be omitted. There is a handful of established kernel functions with different properties. The right choice requires detailed knowledge about the own data and the kernel function, which is why it is recommended to try out different kernel functions.

Discussion SVM

Pro

- **Accuracy:** High classification rate
- **Effective:** Even when number of dimensions $>$ number of samples
- **Robust:** Low tendency to overfitting
- **Compact Models:** “Plane in Space”
- **Versatile:** Applicability of different kernel functions

Contra

- **Efficiency:** Long training phase
- **Complexity:** High implementation effort
- **Black-Box:** Hard to interpret models

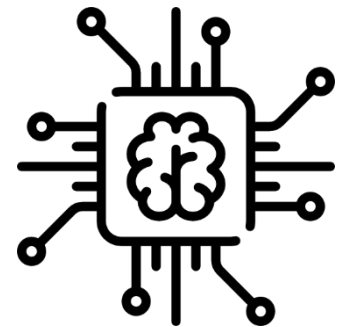
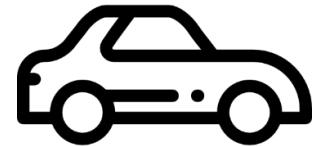
Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

(Andreas Schimpe, M.Sc.)


Agenda

1. Chapter: Introduction
 - 1.1 Overview
 - 1.2 Training and Validation
2. Chapter: Methods
 - 2.1 Logistic Regression
 - 2.2 Nearest Neighbors
 - 2.3 Support Vector Machine
- 3. Chapter: Application**
4. Chapter: Summary



Classification Applications

- Big data
 - Find patterns
 - Make data usable
- Image classification
 - Handwritten Digits
 - X-Rays
- Music classification
 - Shazam
- Speech/Language classification
 - Siri/Alexa/Echo
- Fault detection
 - Quality control during production




Classification for Automotive Technology

- Example: **Perception**

- Camera outputs pixel array
- Classification adds value to each pixel

- Pixel segmentation
- Object detection
- Object tracking



[9]

Vehicle Detection and Tracking



[10]


Vehicle Detection and Tracking

- Get training data
- Extract features from images
- Train classification model based on features
- Take one video frame and classify features of sub-images
- Merge classified areas and create bounding box

Training Data

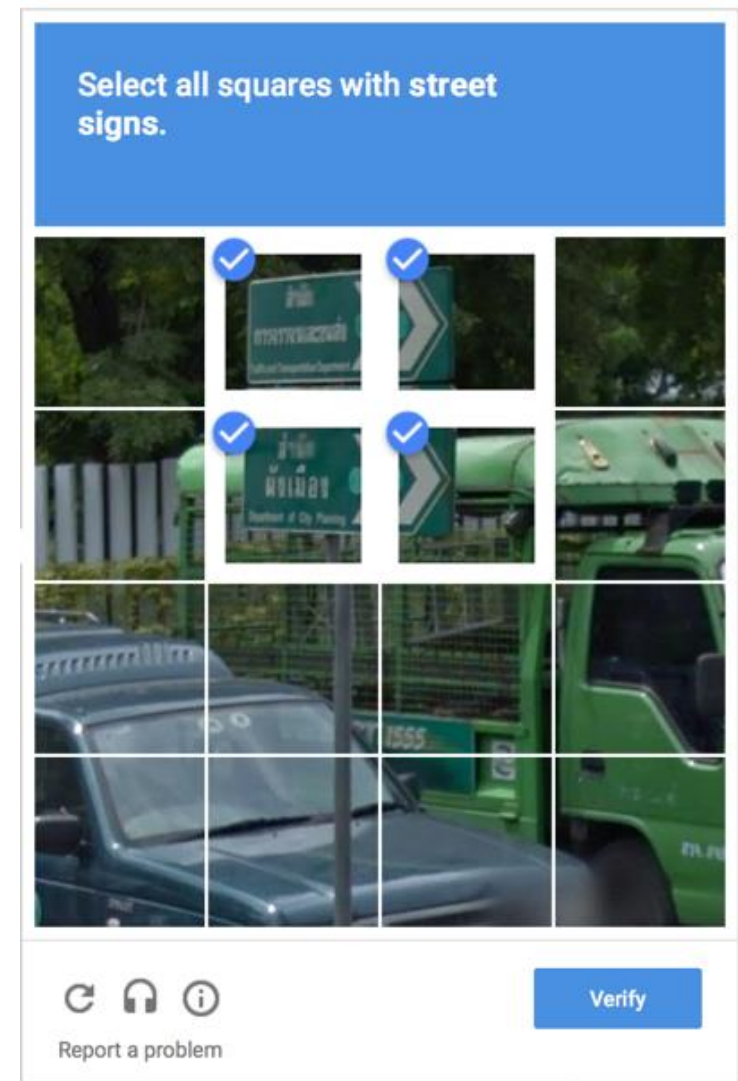
- Required label: „car“ or „no car“
- Required images:
 - Same format used for classification
 - Representative for what is expected to be found in videostream
 - 8000 images (90 % training and 10 % test)

[11]



Training Data

- How to get labeled data?
 - Label data by yourself
 - Pay someone else to label your data
 - Let other label your data for free
- Collection of labeled data
 - Digits: MNIST
 - <http://yann.lecun.com/exdb/mnist/>
 - 70k images
 - Cars: KITTI
 - www.cvlibs.net/datasets/kitti/
 - 80k images



Training Data

„It's terrifying that both of these things are true at the same time in this world:


On the one hand, computers are driving cars.

On the other, the state of the art test to check that you are not a computer is whether you can successfully identify stop signs in pictures.“

- Anonym

Feature Extraction

- Histogram of Oriented Gradients (HOG)
 - compressed & encoded version of the image



[10]

Build SVM Classifier

- Machine learning libraries (python)
 - scikit-learn (<http://scikit-learn.org/>)




```
>>> from sklearn import svm
>>> clf = svm.SVC()
>>> clf.fit(training_features, training_labels)
>>> clf.score(test_features, test_labels)
>>> clf.predict(new_feature)
```

- Training: 1.44 Seconds
- Test: Accuracy = 0.9848
- Prediction: 0 or 1

Classifiy sub-images


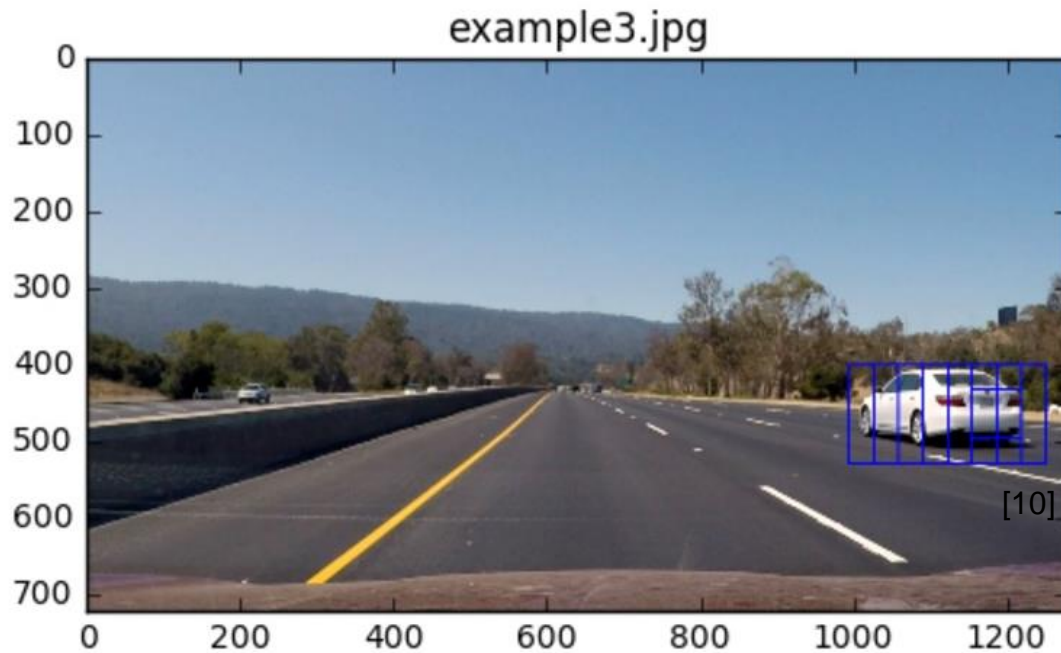
- Produce sub-images of each frame for the classification



[10]

Merge classified areas

- Merge classes of sub-images



Final Output



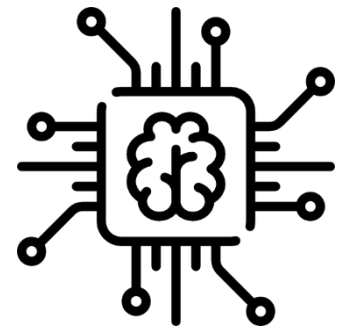
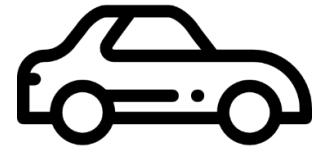
Supervised Learning: Classification

Prof. Dr.-Ing. Markus Lienkamp

(Andreas Schimpe, M.Sc.)

Agenda

1. Chapter: Introduction
 - 1.1 Overview
 - 1.2 Training and Validation
2. Chapter: Methods
 - 2.1 Logistic Regression
 - 2.2 Nearest Neighbors
 - 2.3 Support Vector Machine
3. Chapter: Application
- 4. Chapter: Summary**



Summary

What did we learn today:

- **Classification** is about assigning given classes to data points.
- We need lots of **training data** to build a model for the classification.
- **Machine learning** can extract knowledge from huge datasets.
- Classification is a **supervised learning** problem.
- We need labeled data for training and validation (hidden labels).
- We have several criteria to measure the **quality of a classifier**.
- The concepts of **Logistic regression, Nearest Neighbors and SVM**.
- We can use linear regression together with a sigmoid function as classification method.
- Nearest Neighbors is an instance-based learning method, no training is required.

Summary

What did we learn today:

- **SVMs** are linear classifiers using a maximum margin hyperplane.
- With the **kernel trick**, SVMs can be used for non-linear classification.
- Classification is important for the **perception**, e.g., of cars.
- Acquiring lots of **labeled data** is a problem.
- We have access to good and easy to use **python libraries** for classification.
- We have access to many **open source datasets**, e.g., KITTI for car images.
- Training with big datasets can take **a long time**.
- For image classification, an approach is to **partition, classify (based on features) and then merge** the sub-images.

Sources

- [1] <https://9gag.com/gag/anXDLVn/the-russian-gangster-starter-kit>
- [2] <https://funnyjunk.com/My+neighbours+like+this/funny-pictures/6231925/>
- [3] http://creepypasta.wikia.com/wiki/Alien_Life
- [4] <https://www.youtube.com/watch?v=QopUtQobWJ0>
- [5] <https://www.youtube.com/watch?v=9NrALgHFwTo>
- [6] <https://machinelearningmastery.com/handwritten-digit-recognition-using-convolutional-neural-networks-python-keras/>
- [7] <https://play.google.com/store/apps/details?id=com.shazam.android&hl=de>
- [8] <https://en.wikipedia.org/wiki/Siri>
- [9] <https://playment.io/semantic-segmentation-tool/>
- [10] <https://eu.udacity.com/course/self-driving-car-engineer-nanodegree--nd013>
- [11] www.cvlibs.net/datasets/kitti/
- [12] <https://samsclass.info/120/proj/captchas-021916.htm>
- [13] <https://www.merriam-webster.com/dictionary/classification>
- [14] <https://microage.com/blog/machine-learning-new-cure-all/>
- [15] https://scikit-learn.org/stable/auto_examples/svm/plot_iris.html
- [16] <https://xkcd.com/1838/>

Acknowledgment

- **Machine Learning (Stanford/Coursera)**

- Andrew Ng

- <https://www.coursera.org/learn/machine-learning>

- **Knowledge Discovery in Databases I (LMU)**

- Prof. Dr. Peer Kröger

- http://www.dbs.ifi.lmu.de/cms/studium_lehre/lehre_master/kdd1718/index.html